

Why the Prompts/ Directory Is Genuinely New Ground

A note on what sets a Socratic, clinician-led prompt library apart from every other prompt collection in circulation.

None of what currently exists in the prompt-library space does what the Vibe Rounds **Prompts/** directory is doing. Four things set it apart, and together they mark out territory no other library — clinical or otherwise — has staked out.

1

Socratic Philosophy, Baked In

Most prompt libraries collect prompts that *answer*. This is the only one built around prompts that *question* — backed by a defined pedagogical framework rather than an ad hoc collection of tricks.

2

Domain-Specific Clinical Education Focus

This isn't a generic set of "medical prompts for documentation." It's purpose-built for Socratic clinical reasoning, guided discovery, and teaching ward rounds — a niche that remains otherwise untouched.

3

Validated vs. Unvalidated Distinction

No existing prompt library runs prompts through a human review process for pedagogical soundness. The field relies on community pull requests and star ratings, not expert clinical review.

4

Persona-Based Prompt Architecture

Each prompt embodies a clinical persona — Socratic Attending, Supportive Intern, Junior Resident — with its own constraints and difficulty level attached. That structured, role-calibrated design is novel on its own.

THE SHORT ANSWER

Prompt libraries are common. Medical prompt collections are rare. But a clinician-led, peer-reviewed, Socratic clinical-education prompt library with a published conceptual framework behind it — that doesn't exist anywhere else. The Vibe Rounds Prompts/ directory is genuinely new ground.

What This Asks of You as a Prompt Engineer

Distilled from the project's own working discussions on where Vibe Rounds is headed and what good prompt engineering looks like inside it — a working reference, not a finished rulebook.

Four Standards Every Prompt Must Meet

- Grounded in a real clinical use model or cognitive process — not a generic question template.
- Written by a real domain expert who actually does the work the prompt is meant to support.
- Peer reviewed by another clinician educator before it earns a place in the directory.
- Experimentally validated against real or archived cases, not just reasoned through on paper.

Two Ways to Contribute

Prompt Engineering

Design and refine the personas themselves — Socratic Attending, Supportive Intern, Junior Resident — including the constraints that keep each one questioning instead of answering.

Case-Based Experimentation

Run existing prompts against real or archived cases to see where they hold up, where they break, and what needs tightening before the next version.

The two feed each other: experimentation surfaces what needs fixing in a prompt, and a tightened prompt gives the next experiment something sharper to test against.

The Real Differentiator Is the Prompt Itself

Asking AI to help apply a guideline to a patient is less useful than asking it to evaluate and improve a plan you've already drafted yourself — and more useful still when a learner who's still building that competency drafts the plan and uses Socratic AI feedback to sharpen it. That sequence also carries less risk in practice: it works as a 24/7 reasoning partner with a built-in self-audit, rather than a shortcut around the thinking it's meant to strengthen.

“The differentiator between a shallow use of AI and a deep one is the prompt engineering itself.”

— from the project's design discussions, June 2026

Build Prompts That Fail Safely

1 The AI stays a tool, not an authority

The human remains the decision-maker for anything that affects a real patient. A prompt should never be written in a way that lets that line blur.

2 Train the skill, not just the artifact

Don't hand someone a finished prompt and call it training. Start from a primitive version and let the user watch — or run — it improve step by step, so what transfers is the ability to build prompts, not just one prompt.

3 Let difficulty scale with the user

Lighter friction for self-directed learners; a graded path of increasing rigor — the project's six-level difficulty scale — for everyone else.

Design Against These Biases

Vibe Rounds runs on a real LLM, which means it inherits real failure modes. Part of the prompt engineer's job is designing prompts that account for them, not ignore them.

Question design	Socratic framing that steers reasoning in one direction; over-scaffolding that does the thinking for the learner; under-scaffolding that leaves them stuck; hint dependency.
Model & data	Training-data gaps; recency bias toward newer evidence; historical bias toward older paradigms; language and geographic bias favoring English-language, high-income-country sources.
Evidence selection	Guideline bias toward mainstream recommendations; consensus bias; specialty bias; publication-bias inherited from the literature the model draws on.
User-dependent	The same prompt behaves differently depending on the user's knowledge level, motivation, and how carefully they answer.
Evaluation	Reasoning quality is hard to score objectively, and the cases that go well are the ones anyone remembers — survivorship bias.

“At the highest level, Vibe Rounds is biased by the questions it chooses not to ask — because questions define the search space of reasoning.”

SOURCE

Concept site: avi33tbtt.github.io · Source & docs: github.com/avi33tbtt/Vibe-Rounds

LICENSE

Released under Creative Commons CC BY 4.0. Free to share and adapt with attribution.