

Vibe Rounds — Project Documentation

Author/Concept: Dr. Avinash Kumar Gupta · Coined June 2026 **License:** Creative Commons CC BY 4.0 (free to share with attribution) **Primary sites covered:** - Main site: <https://avi33tbtt.github.io/> - Prompt Modules directory: <https://avi33tbtt.github.io/Prompts/>

Disclaimer carried throughout the project: “To be used on personal responsibility and only for learning. Not to be used for clinical purpose.” All outputs are learning artifacts, not clinical decisions, and require independent verification.

1. What Vibe Rounds Is

Vibe Rounds is a **prompt-based teaching framework**, not a piece of hosted software with persistent state. It is a public library of plain-text “persona” prompts and structured modules designed to make any general-purpose LLM (Claude, ChatGPT, Gemini, etc.) behave like a Socratic clinical attending — questioning a learner’s reasoning, flagging biases, and withholding answers until the learner has committed to a position — rather than an assistant that simply answers on request.

Definition used on the site: > “Vibe Rounds is a Socratic AI paradigm for medical students, clinicians, researchers, and scientists — designed to augment clinical reasoning through structured questioning, guided discovery, and reflective dialogue rather than direct answers.”

The framework positions itself along a **five-stage pedagogical spectrum**, moving from instructor-centered transmission to learner-centered inquiry:

Stage	Name	Educator role	Learner role	Vibe Rounds status
1	Lecture	Expert / Source	Passive receiver	Not a Vibe Rounds focus
2	Dyadic	Facilitator / Partner	Active collaborator	Not a Vibe Rounds focus
3	Socratic	Questioner	Critical thinker	Home base — Socratic Attending persona, answer-withholding constraints
4	Guided Discovery	Architect of experience	Problem solver	Extension — tiered-hint mechanism, Modules/Agents/Pipelines
5	Research	Mentor / Guide	Independent scholar	

Stage	Name	Educator role	Learner role	Vibe Rounds status
				Extension — N-of-1 protocol, evidence mapping

2. Site Structure Overview

The main site (avi33tbtt.github.io) is organized around six main tabs plus a resource/demo library:

Tab	Purpose
Quick Tutorial	Landing page — what it is, learner/educator/researcher quick-start, disclaimer, licensing
Socratic Learning	Persona prompts and behavioral constraints (Stage 3)
Guided Discovery	Modules/Agents/Pipelines architecture, discovery journey, Analytics Mode (Stage 4)
Research	Evidence mapping and the ASV worked demonstration (Stage 5)
More Details	Links to Courseware, EBM Course, Clinical Cognition Course, Project Report, demo apps, YouTube, Terminal, Phronesis, Decision Lab, CAT Bank
Trainings	Live Zoom / on-site training with Dr. Avinash, builder-community access

A separate site section, avi33tbtt.github.io/Prompts/, hosts the full module directory (“VibeRounds Prompt Modules”) — the underlying prompt repository referenced throughout the main site.

3. Stage 3 — Socratic Learning

3.1 Persona Prompts

Short, copy-pasteable `#VibeRounds` commands that a learner pastes as the opening message of a chat session, followed by a clinical scenario. Documented personas include: - **Supportive Intern** — peer-level ward help - **Socratic Attending** — imaginary case generation, or stress-testing a learner’s analysis of a real case - Role-calibrated variants: junior resident, examiner

3.2 Ten Behavioral Constraints

Layered onto any persona prompt to keep the model from lapsing into answer-giving mode:

1. Forced commitment first (no hint until the learner answers)

2. Minimum effort threshold (redirects low-effort replies like “idk”)
3. The 10-second rule (delayed rescue — pause before guiding)
4. Tiered hints, never the full answer up front (framework → narrowed direction → partial answer → summary)
5. Effort-weighted assistance
6. Reasoning grading, not just correctness
7. Adaptive difficulty
8. Role-calibrated rigor
9. Reflection before reveal (“Why do you think that?”, “What could kill the patient?”)
10. Answer-withholding policy (full answer only after genuine attempt or explicit surrender)

3.3 System Prompt Evolution & Difficulty Analytics

A documented, versioned progression (V0–V2) scaling prompt rigor across **six difficulty levels**, from self-directed study (L1) to OSCE-style examination rigor (L6). Available as a downloadable PDF on the main site.

3.4 Reference Implementations

- **Vibe Rounds – Case Bench Lite** — browser-based, bring-your-own-API-key tool. Three-step flow: choose model/provider → paste or upload a case → run a live Socratic exchange. Transcript can be downloaded or anonymously submitted to an in-development research dataset.
- **Level Bench** — same structure as Case Bench Lite, plus a live **L1–L6 level selector** that can be changed mid-conversation without losing the transcript.
- **Socratic Learning Demo App** — standalone live app running the Socratic persona prompts with fixed, non-drifting initialization.

3.5 Maturity

Rated **high maturity** in the framework’s own domain-maturity map: “Socratic feedback loop, six-level difficulty framework, V0–V2 prompt evolution, pilot protocol, and live deployment demonstrated.”

4. Stage 4 — Guided Discovery

4.1 Three-Level Composition Model

Level	Unit	What it does	Examples
1	Modules	A single cognitive lens applied once to a case	Observation, Hypothesis Generation, Bias Detection, Decision Analysis
2	Agents	Orchestrates multiple modules within one run	Guided Discovery Agent, Clinical Cognition Deep Dive, Analytics Agent

Level	Unit	What it does	Examples
3	Pipelines	An explicit, ordered sequence of modules/agents	Example: 1 → 12 → 9 → 21 → 35

Built on the same underlying **57+ module library** and **4 cross-cutting pedagogical frameworks** (see Section 6).

4.2 The Six-Stage Discovery Journey

1. **Observation** — raw case data
2. **Pattern Recognition** — cue clustering
3. **Hypothesis Generation** — differential building
4. **Decision Architecture** — reasoning map
5. **Bias Detection** — cognitive audit (anchoring, premature closure, etc.)
6. **Metacognitive Reflection** — cognition map of the reasoning process itself

Framed as an “MRI for clinical thinking”: instead of asking “What is the diagnosis?”, Guided Discovery asks “How does clinical thinking move from uncertainty to understanding?”

4.3 Analytics Mode — Three Worked Invocation Levels

A minimal recurring pattern: case + module/agent/pipeline reference, submitted as one instruction to any LLM chat interface — no hosting or API integration required.

Invocation	Submitted	Scope	Worked demo case
Single module	Case + Module 42 (Clinical Pre-Mortem)	Focused single-lens pass	42F with severe regular edema
Targeted pipeline	Case + ordered sequence 1 → 12 → 9 → 21 → 35	Progressive multi-stage reasoning	Same case
Full Guided Discovery Agent	Case + agent reference	Deepest single-invocation form; agent internally scores the module catalogue across 8 dimensions	Same case

Each has a linked, unedited Claude.ai transcript for independent inspection.

4.4 Standalone App

Guided Discovery Demo App — live implementation of the seven-stage N-of-1 workflow, with the composition model and discovery journey built into the app’s own fixed scaffolding.

4.5 The N-of-1 Case Research Protocol (bridge to Stage 5)

A **seven-stage N-of-1 case-research workflow**, a two-tier analysis methodology, and outputs aligned to the **CARE case-report guideline**. This is the “guided discovery research” domain in the maturity map — rated **medium maturity**: worked case example complete, multi-case validation still outstanding.

4.6 Maturity & Named Limitations

- No machine-enforced pipeline grammar (any ordered sequence of module IDs is syntactically valid; “analytic yield” is a heuristic, not a validity guarantee)
 - Free-text hand-off between modules (not schema-validated structured data)
 - Single-run demonstrations, not repeated-measures/controlled-comparison evidence
 - Generative variability between runs on identical input
-

5. Stage 5 — Research

5.1 What Clears the Bar

The framework defines “research” (Stage 5) as: **generating or testing something that could hold true beyond the one patient in front of you** — distinct from the single-case reasoning tools (bias audits, SDM prep, critical-thinking audits) that make up most of the Guided Discovery module library.

5.2 Evidence Mapping as the Genre

The adopted secondary-research genre: pulling real published literature on a clinical question and synthesizing **across evidence tiers** (case reports, observational data, systematic reviews, RCTs) **without pooling tiers that shouldn’t be pooled**. This sits one rung below a formal systematic review.

5.3 Worked Demonstration: ASV Mortality-Benefit Evidence Map

What was built: A full evidence-tier synthesis on anti-snake-venom (ASV) mortality benefit, built from PubMed primary literature. Every case report, systematic review, meta-analysis, and RCT surfaced by a documented, filtered PubMed search was extracted individually into a running trajectory map, keeping tiers separate.

- Open demo (uncompressed full network): <https://avi33tbtt.github.io/demo/critical-hub-node-navigation/asv-full.html>
- Compressed entry-point view: <https://avi33tbtt.github.io/demo/critical-hub-node-navigation/index.html>
- Interactive sandbox: <https://avi33tbtt.github.io/demo/critical-hub-node-navigation/asv-simulator.html>

What it found: RCT and meta-analysis tiers do **not** establish an ASV mortality benefit — not because the evidence contradicts one, but because those tiers were never designed/powered to answer that question. The mortality signal lives almost entirely in the **case-report and observational tier**. This is described as a durable, structural feature of the question (a mortality-benefit RCT that withholds antivenom would be unethical to run) rather than a gap awaiting a future trial.

Classification: A scoping/evidence-mapping review — a recognized secondary-research genre, one rung below a systematic review.

Named limitations: - Sampling incomplete — corpus started at 12 of 78 available case reports, being expanded, with an exclusion log - Single-extractor, LLM-assisted extraction — no independent second reviewer yet (spot-checked against source PDFs) - No pre-registered protocol (e.g., PROSPERO) - RCT/meta-analysis and case-report tiers currently analyzed as separate rather than integrated passes

Scope statement: - *What this is:* a transparent, source-traceable way to turn a corpus of published cases into a question or claim that could hold beyond any one patient. - *What this isn't:* a substitute for a registered systematic review, a statistical meta-analysis, or peer review. Every output is a draft to verify against the primary literature, not a citable finding on its own.

6. Clinical Cognition Operating System (CCOS)

CCOS is the architectural frame underlying Modules/Agents/Pipelines — described as “a structured prompt architecture... not software with persistent state or automated checks running outside the chat.”

6.1 Five Generations of Evolution

1. **Gen I** — Prompt library (standalone, answer-oriented, no workflow)
2. **Gen II** — Modular reasoning (illness scripts, differentials, reflective modules)
3. **Gen III** — Protocol networks (dependency-linked cognitive pipelines)
4. **Gen IV** — Metacognitive layer (bias detection, uncertainty mapping, reflection loops)
5. **Gen V — CCOS (current)** — Full clinical cognition OS with trust safeguards; uncertainty-explicit, decision-spectrum, longitudinal learning

6.2 Six Cognitive Layers

1. Clinical reasoning (illness scripts, differential diagnosis, Bayesian inference)
2. Workflow engine (protocol networks, module sequencing, pipeline coordination)
3. Metacognitive monitoring (bias detection, assumption audit, uncertainty mapping)
4. **Epistemic trust layer** — hallucination calibration, prioritized claim verification, confidence calibration (prompt-level framing, not automated fact-checking)
5. Decision architecture (Conservative → Balanced → Maximal spectrum)
6. Learning & documentation (clinical closure, competency tracking, longitudinal archive)

6.3 Four Pedagogical Frameworks (A–D)

Woven into specific steps across modules rather than run standalone:

Framework	Focus
A — Humanistic Persona	Confidence-building; specific affirmation before challenge
B — Fink’s Taxonomy (FLINK)	

Framework	Focus
	Six non-hierarchical learning dimensions (foundational knowledge, application, integration, human dimension, caring, learning-how-to-learn)
C — Bloom’s Revised Taxonomy	Six cognitive levels, Remember → Create, mapped to clinical reasoning tasks
D — Critical Awareness Framework	Standing closing prompt naming the framework’s own susceptible biases (automation bias, anchoring, hallucination risk, rare-diagnosis overweighting)

6.4 Module Lifecycle

Every module follows: **Initiation** → **Execution** → **Closure/Review** — regardless of content, giving each session a consistent shape.

6.5 Scale

- 57+ reasoning modules
- 4 pedagogical frameworks
- 200+ clinical pipelines

6.6 Entry Points by Audience

Audience	Suggested starting modules
Medical students / junior doctors	M01 (Socratic Reasoning), M04 (pre-ward-round prep), M17 (problem representation)
Family caregivers / advocates	M02 (Patient Advocate Documentation), M03 (ongoing monitoring), M11 (plain-language explanations)
Clinicians in practice	M05 (audit a live patient log), M09 (case write-up), M13 (polypharmacy)
Educators / prompt authors	M08 (prompt audit/QC), Framework D (end-of-session bias counter)

Starting point for anyone new: [Module 0 — Cold-Start Orientation](#)

7. Prompt Modules Directory (/Prompts/)

The **VibeRounds Prompt Modules** page documents the underlying repository in more architectural detail, describing it as a layered “learning stack” rather than a flat prompt list:

- **Layer 1 — Pedagogical Frameworks** (A–D, see 6.3): the “foundation layer,” woven into Modules 0–10.

- **Layer 2 — Educational Coverage** (the “content layer”): spans active reasoning (Module 1 — Socratic questioning with answer-withholding), practical rehearsal (Module 4 — ward rounds, admissions, pre-op clearance, polypharmacy, overnight triage), advanced research (Module 9 — seven-stage N-of-1 protocol, the most methodologically formal module), and continuous learning (Module 10 — journal/article reading; Module 7 — registry-based cross-case learning, described as still primitive).
- **Layer 3 — Structured Learning Lifecycle** (the “process layer”): every module runs Initiation → Execution → Closure/Review; Module 1’s closure phase includes a missed-diagnosis debrief and a difficulty-ratchet recommendation for the next session.

The directory notes candidly that **completeness is uneven** — some components are mature and ready for self-directed use, others explicitly experimental — and evaluates the stack layer by layer rather than presenting it as uniformly finished.

Key entry points: - Full prompt directory: <https://avi33tbtt.github.io/Prompts/Prompts.html> - Master Protocol: <https://avi33tbtt.github.io/Prompts/VibeRounds-Master-Protocol.html> - Module 0 (Cold-Start Orientation): <https://avi33tbtt.github.io/Prompts/Module-00-Cold-Start-Orientation.html> - Framework pages: Framework-A/B/C/D-*.html (Humanistic Persona, FLINK, Bloom’s, Critical Awareness)

8. Tools and Applications

8.1 Working Tools (as listed on the Tools Home page)

The Tools Home page is the canonical, numbered directory of Vibe Rounds’ live, browser-based tools. **All run entirely client-side — “bring your own key”: each tool’s API key is stored only in local browser storage and sent directly from the browser to the chosen provider, never through a project server.** These are educational reasoning aids, not diagnostic tools, and not for real patient care.

★ Flagship

#	Tool	Description	Link
★	Tool Hub	All common tools combined into a single workspace/flow — start here if unsure which tool to open	https://avi33tbtt.github.io/hub.html

Core reasoning tools

#	Tool	Description	Link
1	Case Bench	Case-based Socratic learning, MCQs, and clinical reasoning analytics	https://avi33tbtt.github.io/bench.html
2	Case Simulator	Interactive tool for running through simulated patient cases	https://avi33tbtt.github.io/case-simulator.html
3	Clinical Practice		

#	Tool	Description	Link
		Focused drill tool for practicing clinical cases	https://avi33tbtt.github.io/case-practice.html
4	Clinical Polemos	Debate-style exercise for discussion and reasoning	https://avi33tbtt.github.io/clinical-polemos.html
5	Case Lens	Critical thinking around clinical cases from multiple analytical angles	https://avi33tbtt.github.io/case-lens.html
6	SDM Lens	Shared decision-making scenario tool	https://avi33tbtt.github.io/sdm-lens.html
7	EBM Query Generator	Generates a structured set of critical-appraisal queries	https://avi33tbtt.github.io/ebm.html
8	Clinical Cognition OS (CCOS) — Module Order Builder	Build and order modules for CCOS; outputs a ready-to-copy Analytics Mode prompt	https://avi33tbtt.github.io/ccos-builder.html
9	Case Analytics	Analytics tool for reviewing clinical reasoning performance	https://avi33tbtt.github.io/case-analytics.html
10	Case Reading	Guided reading and analysis of clinical cases	https://avi33tbtt.github.io/case-reading.html
11	Critical Hub-Node Navigation	“The Avinash Principle” — find the handful of nodes that decide the outcome	https://avi33tbtt.github.io/critical-hub-node-navigation.html
12	Complex Flow Builder	Maps and builds out complex clinical decision flows (replace case placeholder with your own case data/link)	https://avi33tbtt.github.io/complex-flow-builder.html
13	CQM Router	Clinical query/module router — directs cases to the right tool or pathway	https://avi33tbtt.github.io/cqm-router.html
14	Case Question Bank	A bank of clinical case questions for practice and review	https://avi33tbtt.github.io/case-question-bank.html

Courses

#	Tool	Description	Link
15	Courseware	Structured course material and modules for guided clinical learning (19 subjects, 190 chapters, built-in Prompt Builder)	https://avi33tbtt.github.io/courseware/

#	Tool	Description	Link
16	EBM	Course on evidence-based medicine — critically appraising and applying clinical evidence	https://avi33tbtt.github.io/ebm/
17	Clinical Cognition	Course content on critical care reasoning and decision-making	https://avi33tbtt.github.io/cc/

Tutor, Article, Research

#	Tool	Description	Link
18	Tutor	Interactive, drill-style tools for practicing clinical reasoning	https://avi33tbtt.github.io/#tutor-panel
19	Article	Background reading on the Vibe Rounds paradigm and its approach to clinical reasoning	https://avi33tbtt.github.io/article.html
20	Research (Bench Lite)	Run a Socratic learning session, then hit Submit at the end to help gather data for research	https://avi33tbtt.github.io/bench_lite.html

Pro tip from the site: upload a tool's text output into any LLM and ask it to convert that into an HTML or PDF file to save as personal notes.

8.2 Additional Tools and Demo Apps (from the main site and elsewhere)

Tool	Description	Link
Case Bench Lite	Lightweight single-difficulty Socratic session (same as #20 above)	https://avi33tbtt.github.io/bench_lite.html
Level Bench	Same structure as Case Bench Lite, plus a live L1–L6 difficulty control, adjustable mid-session	https://avi33tbtt.github.io/level_bench.html
Socratic Learning Demo App	Standalone Stage-3 live app with fixed, non-drifting initialization	https://viberoundslearning.ai.studio/
Guided Discovery Demo App	Standalone Stage-4 live app (7-stage N-of-1 workflow)	https://viberoundsdiscovery.ai.studio/
Terminal	Command-line console for direct low-level interaction with CCOS primitives	https://avi33tbtt.github.io/terminal/
Phronesis	Interactive practical-wisdom/decision-spine module	https://avi33tbtt.github.io/phronesis/

Tool	Description	Link
Decision Lab	Commitment-gated branching walkthrough (diagnosis/therapy/prognosis/screening)	https://avi33tbtt.github.io/Decision-Lab/anemia.html
CAT Bank	30 critically appraised topics on anaemia	https://avi33tbtt.github.io/CATBank/anemia.html
System 3 — Ultima Thule	Meta-cognitive “overseer” layer auditing systemic bias across multi-module setups	https://avi33tbtt.github.io/system3/
Illness Velocity	Graph-based rate-of-change dashboard for anticipating acute decompensation	https://avi33tbtt.github.io/demo/Velocity/index.html
Agent Demo	Multiple specialized AI personas cross-examining each other in real time	https://avi33tbtt.github.io/demo/agent/index.html

9. Demo Case Library (Selected)

Organized by stage of clinical reasoning, from foundational recall to systems-level synthesis:

- **Flagship:** 14M Microcytic Anemia & Urinary Obstruction Dashboard, VibeRounds Case Reader
- **Foundational:** EBM Demo, EBM Queries, Question Bank Demo
- **Application:** Case-Based MCQ, Patient Education Demo, Clinical Pathways — acute pancreatitis
- **Analysis:** Case Lens 8-Lens Audit, 6-Case Nested Analysis (TMNG)
- **Evaluation:** Decision Lab, SDM Lens, PaJR Analytics
- **Synthesis:** Case Simulator, Phronesis Case 1, Pipelines Demo, Module CC — Post-COVID Hydropneumothorax, VibeRounds Case Demo
- **Systems-level / Research:** ASV — Critical Hub-Node Navigation, ASV Full, ASV Simulator

10. Domain Maturity Map (as self-reported, June/July 2026)

Domain	Maturity	Notes
1. AI-augmented clinical education (Socratic Learning)	High	Socratic feedback loop, six-level difficulty framework, V0–V2 prompt evolution, pilot protocol, live deployment demonstrated
2. Guided discovery research (N-of-1 protocol, Research layer)	Medium	Seven-stage N-of-1 workflow, two-tier analysis, CARE-aligned outputs defined; worked case complete; multi-case validation outstanding
3. Bedside clinical decision support	Early stage	

Domain	Maturity	Notes
		Concept/architecture defined; EMR integration, multimodal AI layer, FHIR infrastructure pre-implementation

11. Related Academic Publications

All hosted open-access on ResearchGate (profile: <https://www.researchgate.net/profile/Avinash-Kumar-Gupta-2>):

1. *Socratic Learning in Medical Education* — https://www.researchgate.net/publication/411022653_Socratic_Learning_in_Medical_Education
 2. *Guided Discovery in Medical Education: Scaffolded Environments, Modules, Agents, and Pipelines in the Vibe Rounds CCOS Framework* — <https://www.researchgate.net/publication/411022105>
 3. *Research in the Vibe Rounds Framework: Evidence Mapping, Stage 5, and a Worked Demonstration on Anti-Snake-Venom Mortality Benefit* — <https://www.researchgate.net/publication/411031621>
 4. *Modern Clinical Reasoning & Socratic AI: Integrating Prompt and Provocation Cognitive Modes into Healthcare Analytics* — <https://avi33tbtt.github.io/article.html>
 5. *Prompt Architecture as a Teaching Modality: A Commentary for Medical Educators on the Vibe Rounds Framework — Every Layer and Two Worked Case Studies* — <https://www.researchgate.net/publication/411102441>
 6. *From Prompt Library to Operating System: The Clinical Cognition Operating System (CCOS) as the Architectural Layer Beneath the Vibe Rounds Framework* — <https://www.researchgate.net/publication/411101664>
-

12. Contact and Training

- Email: avi33tbtt@gmail.com
 - WhatsApp: +977 981 759 9973 / +91 731 851 0434
 - LinkedIn: <https://www.linkedin.com/in/dravinashkumargupta/>
 - YouTube: https://www.youtube.com/@geeky_dr
 - Training options: online Zoom small-group sessions, on-site departmental training; both include access to an ongoing builder community of educators/clinicians/researchers extending the framework.
-

13. Overall Framing and Caveats

Across every layer, the project consistently self-labels its own evidentiary status rather than overstating it:

- It is a **prompt architecture**, not software with persistent state, automated verification, or a machine-enforced pipeline grammar.
- Outputs are **generative and non-deterministic** — an identical case and module sequence can produce different output on repeated runs.
- Every worked demonstration (Analytics Mode examples, ASV evidence map, case dashboards) is presented as an **illustration of the pattern working as designed**, not as validated outcome evidence or a citable finding on its own.
- The persistent site-wide disclaimer: **not for clinical use**, learning purposes only, all outputs require independent clinical verification.

This documentation was compiled from the live content of <https://avi33tbtt.github.io/> and <https://avi33tbtt.github.io/Prompts/> as of August 2026, together with the project's own companion papers on the Socratic Learning, Guided Discovery, and Research layers. Site content may change; consult the live pages for the current state of the project.